

Branch: **master** ▼

Find file

Copy path

[paxosstore](#) / [paxoskv](#) / [core](#) / [pins_wrapper.cc](#)**dengoswei** - pass pins_wrapper_test;

1307944 on Aug 27, 2017

[1 contributor](#)

Raw

Blame

History



966 lines (838 sloc) 29.8 KB

```
1
2  /*
3   * Tencent is pleased to support the open source community by making PaxosStore available.
4   * Copyright (C) 2017 THL A29 Limited, a Tencent company. All rights reserved.
5   * Licensed under the BSD 3-Clause License (the "License"); you may not use this file
6   * without permission. You may obtain a copy of the License at
7   * https://opensource.org/licenses/BSD-3-Clause
8   * Unless required by applicable law or agreed to in writing, software distributed under the
9   * License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
10  * KIND, either express or implied. See the License for the specific language governing
11  * permissions and limitations under the License.
12
13  #include <unistd.h>
14  #include <cassert>
15  #include "pins_wrapper.h"
16  #include "cutils/mem_utils.h"
17  #include "cutils/log_utils.h"
18  #include "cutils/hassert.h"
19
20  namespace {
21
22  inline void set_accepted_value(
23      std::unique_ptr<paxos::Message>& rsp_msg,
24      const paxos::Entry& accepted_value)
25  {
26      assert(nullptr != rsp_msg);
27      auto entry = rsp_msg->mutable_accepted_value();
28      assert(nullptr != entry);
29      *entry = accepted_value;
30      assert(rsp_msg->accepted_value().reqid() == accepted_value.reqid());
31      assert(rsp_msg->accepted_value().data() == accepted_value.data());
32  }
33
34  void updateRspVotes(
```

```
35     uint8_t peer_id,
36     bool vote,
37     std::map<uint8_t, bool>& rsp_votes)
38 {
39     assert(0 < peer_id);
40
41     if (rsp_votes.end() != rsp_votes.find(peer_id)) {
42         assert(rsp_votes[peer_id] == vote);
43         return ;
44     }
45
46     // else
47     rsp_votes[peer_id] = vote;
48 }
49
50 std::tuple<int, int> countVotes(const std::map<uint8_t, bool>& votes)
51 {
52     int true_cnt = 0;
53     int false_cnt = 0;
54     for (const auto& v : votes) {
55         if (v.second) {
56             ++true_cnt;
57         } else {
58             ++false_cnt;
59         }
60     }
61
62     return std::make_tuple(true_cnt, false_cnt);
63 }
64
65 inline bool
66 updatePromised(uint64_t prop_num, paxos::PaxosInstance& pins_impl)
67 {
68     if (pins_impl.has_promised_num() &&
69         pins_impl.promised_num() > prop_num) {
70         return false; // reject
71     }
72
73     pins_impl.set_promised_num(prop_num);
74     return true;
75 }
76
77 bool updateAccepted(
78     uint64_t prop_num,
79     const paxos::Entry& prop_value,
80     bool is_fast_accept,
81     paxos::PaxosInstance& pins_impl)
82 {
```

```
83     if (pins_impl.has_promised_num() &&
84         pins_impl.promised_num() > prop_num) {
85         return false; // reject
86     }
87
88     assert(false == pins_impl.has_promised_num() ||
89         pins_impl.promised_num() <= prop_num);
90     if (true == is_fast_accept) {
91         if (pins_impl.has_accepted_num()) {
92             // do fast accepted only when accepted_num haven't be set
93             return false; // reject
94         }
95
96         assert(false == pins_impl.has_accepted_num());
97     }
98
99     assert(false == pins_impl.has_accepted_num() ||
100         pins_impl.accepted_num() <= prop_num);
101     pins_impl.set_promised_num(prop_num);
102     pins_impl.set_accepted_num(prop_num);
103     {
104         auto entry = pins_impl.mutable_accepted_value();
105         assert(nullptr != entry);
106         *entry = prop_value;
107         assert(pins_impl.accepted_value().reqid() == prop_value.reqid());
108         assert(pins_impl.accepted_value().data() == prop_value.data());
109     }
110
111     return true;
112 }
113
114
115 } // namespace
116
117 namespace paxos {
118
119 // function for test
120
121 std::unique_ptr<PInsAliveState> PInsAliveState::TestClone()
122 {
123     auto clone_pins_state =
124         cutils::make_unique<PInsAliveState>(key_, index_, prop_num_gen_.Get()
125     assert(nullptr != clone_pins_state);
126     assert(clone_pins_state->prop_num_gen_.Get() == prop_num_gen_.Get());
127
128     clone_pins_state->prop_state_ = prop_state_;
129     clone_pins_state->max_accepted_hint_num_ = max_accepted_hint_num_;
130     clone_pins_state->max_hint_num_ = max_hint_num_;
```

```
131     clone_pins_state->rsp_votes_ = rsp_votes_;
132     if (nullptr != proposing_value_)
133     {
134         clone_pins_state->proposing_value_ =
135             cutils::make_unique<Entry>(*proposing_value_);
136         assert(nullptr != clone_pins_state->proposing_value_);
137     }
138
139     return clone_pins_state;
140 }
141
142 // end of function for test
143
144 PinsAliveState::PinsAliveState(
145     const std::string& key,
146     uint64_t index,
147     uint64_t proposed_num)
148     : key_(key)
149     , index_(index)
150     , prop_num_gen_(proposed_num)
151 {
152     assert(0 < index);
153     assert(0 == pipe(pipes_));
154     assert(0 <= pipes_[0]);
155     assert(0 <= pipes_[1]);
156 }
157
158 PinsAliveState::~PinsAliveState()
159 {
160     assert(0 <= pipes_[0]);
161     assert(0 <= pipes_[1]);
162     close(pipes_[1]);
163     close(pipes_[0]);
164     pipes_[0] = -1;
165     pipes_[1] = -1;
166 }
167
168
169 void PinsAliveState::MarkChosen()
170 {
171     prop_state_ = PropState::CHOSEN;
172     rsp_votes_.clear();
173     proposing_value_ = nullptr;
174     assert(IsChosen());
175 }
176
177 void PinsAliveState::SendNotify() const
178 {
```

```

179     assert(0 <= pipes_[0]);
180     assert(0 <= pipes_[1]);
181     char ch = 'p';
182     assert(1 == write(pipes_[1], &ch, 1));
183 }
184
185
186 PropState
187 PinsAliveState::stepPrepareRsp(
188     uint8_t peer_id,
189     uint64_t peer_promised_num,
190     uint64_t peer_accepted_num,
191     const paxos::Entry* peer_accepted_value)
192 {
193     assert(0 < peer_id);
194     assert(PropState::WAIT_PREPARE == prop_state_);
195
196     assert(nullptr != proposing_value_);
197     uint64_t proposed_num = prop_num_gen_.Get();
198     // CHECK proposed_num == peer_promised_num
199     assert(proposed_num <= peer_promised_num);
200     updateRspVotes(
201         peer_id, proposed_num == peer_promised_num, rsp_votes_);
202
203     if (nullptr != peer_accepted_value) {
204         assert(0 < peer_accepted_num);
205         if (peer_accepted_num >= max_accepted_hint_num_) {
206             if (peer_accepted_num > max_accepted_hint_num_) {
207                 max_accepted_hint_num_ = peer_accepted_num;
208                 *proposing_value_ = *peer_accepted_value;
209             }
210             else {
211                 assert(proposing_value_->reqid() == peer_accepted_v
212                 assert(proposing_value_->data() == peer_accepted_va
213             }
214         }
215     }
216
217     // else => reject
218     int promised_cnt = 0;
219     int reject_cnt = 0;
220     std::tie(promised_cnt, reject_cnt) = countVotes(rsp_votes_);
221     if (reject_cnt >= major_cnt_) {
222         // reject by majority
223         return PropState::PREPARE;
224     }
225     else if (promised_cnt + 1 >= major_cnt_) {
226         // +1 => including self-vote

```

```
227         return PropState::ACCEPT;
228     }
229
230     return PropState::WAIT_PREPARE;
231 }
232
233 PropState
234 PinsAliveState::stepAcceptRsp(
235     uint8_t peer_id,
236     uint64_t peer_accepted_num,
237     bool is_fast_accept_rsp)
238 {
239     assert(0 < peer_id);
240     assert(PropState::WAIT_ACCEPT == prop_state_);
241
242     uint64_t proposed_num = prop_num_gen_.Get();
243     updateRspVotes(peer_id, proposed_num == peer_accepted_num, rsp_votes_);
244
245     int accept_cnt = 0;
246     int reject_cnt = 0;
247     std::tie(accept_cnt, reject_cnt) = countVotes(rsp_votes_);
248     if (reject_cnt >= major_cnt_) {
249         return PropState::PREPARE;
250     }
251     else if (accept_cnt + 1 >= major_cnt_) {
252         return PropState::CHOSEN;
253     }
254
255     return PropState::WAIT_ACCEPT;
256 }
257
258 PropState PinsAliveState::stepTryPropose(
259     uint64_t hint_proposed_num,
260     const paxos::Entry& try_proposing_value)
261 {
262     // delete prop_state_ = PropState::PREPARE;
263     prop_num_gen_.Update(hint_proposed_num);
264     active_prop_cnt_ = 0;
265     max_accepted_hint_num_ = 0;
266     if (nullptr == proposing_value_) {
267         proposing_value_ = cutils::make_unique<Entry>(try_proposing_value);
268     }
269     else {
270         *proposing_value_ = try_proposing_value;
271     }
272
273     // else => nothing change..
274     assert(nullptr != proposing_value_);
```

```
275
276     rsp_votes_.clear();
277     return PropState::PREPARE;
278 }
279
280 PropState PinsAliveState::stepBeginPropose(
281     uint64_t hint_proposed_num,
282     const paxos::Entry& proposing_value)
283 {
284     assert(PropState::NIL == prop_state_);
285     prop_num_gen_.Update(hint_proposed_num);
286     // delete prop_state_ = PropState::PREPARE;
287     assert(nullptr == proposing_value_);
288     proposing_value_ = cutils::make_unique<Entry>(proposing_value);
289     return PropState::PREPARE;
290 }
291
292 PropState PinsAliveState::beginPreparePhase(PaxosInstance& pins_impl)
293 {
294     assert(PropState::PREPARE == prop_state_);
295     assert(nullptr != proposing_value_);
296
297     pins_impl.set_proposed_num(prop_num_gen_.Get());
298     if (pins_impl.has_accepted_num() &&
299         max_accepted_hint_num_ < pins_impl.accepted_num()) {
300         assert(pins_impl.has_accepted_value());
301         max_accepted_hint_num_ = pins_impl.accepted_num();
302         *proposing_value_ = pins_impl.accepted_value();
303     }
304
305     if (false == updatePromised(prop_num_gen_.Get(), pins_impl)) {
306         // reject
307         return PropState::PREPARE;
308     }
309
310     rsp_votes_.clear();
311     return PropState::WAIT_PREPARE;
312 }
313
314 PropState PinsAliveState::beginAcceptPhase(PaxosInstance& pins_impl)
315 {
316     assert(PropState::ACCEPT == prop_state_);
317     assert(nullptr != proposing_value_);
318     if (false == updateAccepted(
319         prop_num_gen_.Get(), *proposing_value_, false, pins_impl)) {
320         // reject
321         return PropState::PREPARE;
322     }
323 }
```

```
323
324     // reject promised may bring max_accepted_hint_num_ > pins_impl.accepted_num_
325     // pins_impl.accepted_num may < max_accepted_hint_num_;
326     // assert(pins_impl.accepted_num() >= max_accepted_hint_num_);
327
328     rsp_votes_.clear();
329     return PropState::WAIT_ACCEPT;
330 }
331
332
333
334 std::tuple<bool, MessageType>
335 PinsAliveState::updatePropState(
336     PropState next_prop_state, PaxosInstance& pins_impl)
337 {
338     bool write = false;
339     auto rsp_msg_type = MessageType::NOOP;
340     prop_state_ = next_prop_state;
341     switch (prop_state_) {
342     case PropState::PROP_FROZEN:
343         logerr("REACHE MAX_PROP_CNT %d => PROP_FROZEN", MAX_PROP_CNT);
344         assert(false == write);
345         SendNotify(); // it's safe here!!
346         break;
347
348     case PropState::WAIT_PREPARE:
349     case PropState::WAIT_ACCEPT:
350         // nothing
351         break;
352
353     case PropState::CHOSEN:
354         rsp_msg_type = MessageType::CHOSEN;
355         rsp_votes_.clear();
356         proposing_value_ = nullptr;
357         break;
358
359     case PropState::PREPARE:
360     {
361         ++active_prop_cnt_;
362         if (active_prop_cnt_ > MAX_PROP_CNT)
363         {
364             // MAX_PROP_CNT reached !!
365             return updatePropState(PropState::PROP_FROZEN, pins_impl);
366         }
367
368         if (pins_impl.has_promised_num()) {
369             prop_num_gen_.Next(pins_impl.promised_num());
370         }
```



```
371
372     prop_num_gen_.Update(
373         std::max(max_hint_num_, max_accepted_hint_n
374     hassert(prop_num_gen_.Get() > pins_impl.proposed_num(),
375         "prop_num_gen_.Get %" PRIu64
376         " pins_impl.proposed_num %" PRIu64,
377         prop_num_gen_.Get(), pins_impl.proposed_num());
378     assert(prop_num_gen_.Get() > pins_impl.promised_num());
379     auto new_state = beginPreparePhase(pins_impl);
380     assert(PropState::WAIT_PREPARE == new_state);
381
382     bool new_write = false;
383     MessageType tmp_rsp_msg_type = MessageType::NOOP;
384     std::tie(new_write, tmp_rsp_msg_type)
385         = updatePropState(PropState::WAIT_PREPARE, pins_impl);
386     assert(false == new_write);
387     assert(MessageType::NOOP == tmp_rsp_msg_type);
388     assert(PropState::WAIT_PREPARE == prop_state_);
389     rsp_msg_type = MessageType::PROP;
390     write = true;
391 }
392 break;
393
394 case PropState::ACCEPT:
395 {
396     auto new_state = beginAcceptPhase(pins_impl);
397     if (PropState::PREPARE == new_state) {
398         return updatePropState(PropState::PREPARE, pins_impl);
399     }
400
401     assert(PropState::WAIT_ACCEPT == new_state);
402
403     bool new_write = false;
404     MessageType tmp_rsp_msg_type = MessageType::NOOP;
405     std::tie(new_write, tmp_rsp_msg_type)
406         = updatePropState(PropState::WAIT_ACCEPT, pins_impl);
407     assert(false == new_write);
408     assert(MessageType::NOOP == tmp_rsp_msg_type);
409     assert(PropState::WAIT_ACCEPT == prop_state_);
410     rsp_msg_type = MessageType::ACCPT;
411     write = true;
412 }
413 break;
414
415 default:
416     assert(false);
417 }
418
```

```

419     return std::make_tuple(write, rsp_msg_type);
420 }
421
422
423 std::tuple<bool, MessageType>
424 PinsAliveState::Step(const Message& msg, PaxosInstance& pins_impl)
425 {
426     assert(key_ == msg.key());
427     assert(index_ == msg.index());
428     assert(PropState::CHOSEN != prop_state_);
429
430     bool write = false;
431     MessageType rsp_msg_type = MessageType::NOOP;
432     switch (msg.type()) {
433     case MessageType::PROP_RSP:
434     {
435         assert(nullptr != proposing_value_);
436         if (PropState::WAIT_PREPARE != prop_state_ ||
437             pins_impl.proposed_num() != msg.proposed_num()) {
438             logdebug("msgtype::PROP_RSP "
439                 " index %" PRIu64
440                 " but ins in state %d"
441                 " pins_impl.proposed_num %" PRIu64
442                 " msg.proposed_num %" PRIu64,
443                 msg.index(),
444                 static_cast<int>(prop_state_),
445                 pins_impl.proposed_num(),
446                 msg.proposed_num());
447             break;
448         }
449
450         assert(PropState::WAIT_PREPARE == prop_state_);
451         PropState next_prop_state = PropState::NIL;
452         if (prop_num_gen_.Get() != msg.proposed_num())
453         {
454             // must be write failed
455             assert(prop_num_gen_.Get() > msg.proposed_num());
456             next_prop_state = PropState::PREPARE; // redo
457         }
458         else
459         {
460             assert(prop_num_gen_.Get() == msg.proposed_num());
461             assert(prop_num_gen_.Get() == pins_impl.proposed_num());
462             next_prop_state = stepPrepareRsp(
463                 msg.from(), msg.promised_num(), msg
464                 msg.has_accepted_value()
465                 ? &msg.accepted_value() : n
466             }

```

```

467
468         std::tie(write, rsp_msg_type)
469         = updatePropState(next_prop_state, pins_impl);
470     // valid check
471     {
472         if (MessageType::ACCP_T == rsp_msg_type) {
473             assert(msg.proposed_num() == pins_impl.proposed_num());
474             assert(msg.proposed_num() == pins_impl.promised_num());
475             assert(msg.proposed_num() == pins_impl.accepted_num());
476         }
477     }
478 }
479 break;
480 case MessageType::ACCP_T_RSP:
481 case MessageType::FAST_ACCP_T_RSP:
482 {
483     assert(nullptr != proposing_value_);
484     if (PropState::WAIT_ACCEPT != prop_state_
485         || pins_impl.proposed_num() != msg.proposed_num()) {
486         logdebug("msg ACCP_T_RSP index %" PRIu64
487                 " but instance in state %d"
488                 " pins_impl.proposed_num %" PRIu64
489                 " msg.proposed_num %" PRIu64,
490                 msg.index(),
491                 static_cast<int>(prop_state_),
492                 pins_impl.proposed_num(),
493                 msg.proposed_num());
494         break;
495     }
496
497     assert(PropState::WAIT_ACCEPT == prop_state_);
498     assert(prop_num_gen_.Get() == msg.proposed_num());
499     assert(prop_num_gen_.Get() == pins_impl.proposed_num());
500     assert(msg.has_accepted_num());
501     assert(false == msg.has_accepted_value());
502     auto next_prop_state = stepAcceptRsp(
503         msg.from(), msg.accepted_num(),
504         MessageType::FAST_ACCP_T_RSP == msg.type());
505
506     // valid check
507     {
508         if (PropState::CHOSEN == next_prop_state) {
509             assert(msg.proposed_num() == pins_impl.proposed_num());
510             assert(msg.proposed_num() <= pins_impl.promised_num());
511             assert(msg.proposed_num() <= pins_impl.accepted_num());
512             // MUST BE:
513             // event if pins_impl.accepted_num > msg.proposed_num
514             hassert(proposing_value_>reqid() ==

```

```
515         pins_impl.accepted_value().reqid(),
516         "proposing_value_->reqid %"
517         " pins_impl.accepted_value(
518         proposing_value_->reqid(),
519         pins_impl.accepted_value().
520         assert(proposing_value_->data() ==
521         pins_impl.accepted_value().data());
522     }
523 }
524 std::tie(write, rsp_msg_type)
525     = updatePropState(next_prop_state, pins_impl);
526
527 // update max_hint_num_
528 if (msg.has_promised_num()) {
529     max_hint_num_ = std::max(max_hint_num_, msg.promised_num());
530 }
531
532 if (msg.has_accepted_num()) {
533     max_hint_num_ = std::max(max_hint_num_, msg.accepted_num());
534 }
535 }
536 break;
537 case MessageType::TRY_REDO_PROP:
538 case MessageType::TRY_PROP:
539 {
540     PropState next_prop_state = PropState::NIL;
541
542     uint64_t hint_proposed_num = msg.proposed_num();
543     if (0 == cutils::get_prop_cnt(hint_proposed_num)) {
544         hint_proposed_num = cutils::prop_num_compose(0, 1);
545     }
546
547     assert(msg.has_accepted_value());
548     next_prop_state = stepTryPropose(
549         hint_proposed_num, msg.accepted_value());
550     assert(PropState::PREPARE == next_prop_state);
551
552     assert(0 == active_prop_cnt_);
553     assert(0 == max_accepted_hint_num_);
554
555     std::tie(write, rsp_msg_type)
556         = updatePropState(next_prop_state, pins_impl);
557     assert(PropState::WAIT_PREPARE == prop_state_);
558     assert(rsp_votes_.empty());
559
560     active_begin_prop_num_ = pins_impl.proposed_num();
561 }
562 break;
```

```
563     case MessageType::BEGIN_PROP:
564     case MessageType::BEGIN_FAST_PROP:
565         {
566             // assert(0 == msg.proposed_num());
567             assert(nullptr == proposing_value_);
568             assert(0 == active_prop_cnt_);
569             assert(0 == max_accepted_hint_num_);
570             // use msg.accepted_value as propose value
571             assert(msg.has_accepted_value());
572             if (pins_impl.has_promised_num()) {
573                 logerr("CONFLICT");
574                 break;
575             }
576
577             assert(PropState::NIL == prop_state_);
578             assert(false == pins_impl.has_promised_num());
579             assert(false == pins_impl.has_accepted_num());
580             assert(false == pins_impl.has_accepted_value());
581
582             uint64_t hint_proposed_num = 0;
583             if (MessageType::BEGIN_PROP == msg.type()) {
584                 hint_proposed_num = cutils::prop_num_compose(0, 1);
585             }
586
587             auto next_prop_state = stepBeginPropose(
588                 hint_proposed_num, msg.accepted_value());
589             assert(PropState::PREPARE == next_prop_state);
590             std::tie(write, rsp_msg_type)
591                 = updatePropState(next_prop_state, pins_impl);
592
593             assert(true == write);
594             assert(MessageType::PROP == rsp_msg_type);
595             hassert(prop_num_gen_.Get() == pins_impl.proposed_num(),
596                 "prop_num_gen_.Get %" PRIu64
597                 " pins_impl.proposed_num %" PRIu64,
598                 prop_num_gen_.Get(), pins_impl.proposed_num());
599             assert(pins_impl.has_promised_num());
600             assert(prop_num_gen_.Get() == pins_impl.promised_num());
601             assert(false == pins_impl.has_accepted_num());
602             assert(false == pins_impl.has_accepted_value());
603
604             if (MessageType::BEGIN_FAST_PROP == msg.type()) {
605                 // fast prop
606                 // => skip prepare phase
607                 std::tie(write, rsp_msg_type)
608                     = updatePropState(PropState::ACCEPT, pins_impl);
609                 assert(true == write);
610                 assert(MessageType::ACCPT == rsp_msg_type);
```

```
611         assert(pins_impl.has_accepted_num());
612         assert(pins_impl.has_accepted_value());
613         rsp_msg_type = MessageType::FAST_ACCPT;
614         assert(0 == cutils::get_prop_cnt(pins_impl.proposed
615     }
616
617         assert(0 == max_accepted_hint_num_);
618         assert(rsp_votes_.empty());
619         active_begin_prop_num_ = pins_impl.proposed_num();
620     }
621     break;
622 default:
623     assert(false);
624 }
625
626 return std::make_tuple(write, rsp_msg_type);
627 }
628
629 PinsWrapper::PinsWrapper(
630     PinsAliveState* pins_state,
631     PaxosInstance& pins_impl)
632 : pins_state_(pins_state)
633 , pins_impl_(pins_impl)
634 {
635     if (pins_impl.chosen()) {
636         assert(pins_impl.has_accepted_value());
637     }
638 }
639
640 std::tuple<int, bool, std::unique_ptr<Message>>
641 PinsWrapper::Step(const Message& msg)
642 {
643     assert(msg.index() == pins_impl_.index());
644     if (IsChosen()) {
645         return stepChosen(msg);
646     }
647
648     assert(false == IsChosen());
649     return stepNotChosen(msg);
650 }
651
652 std::tuple<int, bool, std::unique_ptr<Message>>
653 PinsWrapper::stepChosen(const Message& msg)
654 {
655     assert(true == IsChosen());
656     bool write = false;
657     std::unique_ptr<Message> rsp_msg = nullptr;
658
```

```

659     assert(true == pins_impl_.has_promised_num());
660     assert(true == pins_impl_.has_accepted_num());
661     assert(true == pins_impl_.has_accepted_value());
662     switch (msg.type()) {
663     case MessageType::CHOSEN:
664         // check
665         if (msg.has_accepted_value())
666         {
667             // TODO:
668                 if ((msg.accepted_value().data() !=
669                     pins_impl_.accepted_value().data()) ||
670                     (msg.accepted_value().reqid() !=
671                     pins_impl_.accepted_value().reqid()))
672                 {
673                     logerr("IMPORTANT INCONSISTENT index %" PRIu64
674                             " from %u to %u",
675                             msg.index(), msg.from(), msg.to());
676                     assert(false);
677                 }
678             }
679     default:
680         break;
681
682     case MessageType::GET_CHOSEN:
683     case MessageType::PROP:
684     case MessageType::ACCP:
685     case MessageType::FAST_ACCP:
686         rsp_msg = cutils::make_unique<Message>();
687         assert(nullptr != rsp_msg);
688
689         rsp_msg->set_type(MessageType::CHOSEN);
690         rsp_msg->set_index(msg.index());
691         rsp_msg->set_key(msg.key());
692         rsp_msg->set_from(msg.to());
693         rsp_msg->set_to(msg.from());
694
695         rsp_msg->set_proposed_num(pins_impl_.proposed_num());
696         rsp_msg->set_promised_num(pins_impl_.promised_num());
697         rsp_msg->set_accepted_num(pins_impl_.accepted_num());
698         rsp_msg->set_timestamp(time(NULL));
699         set_accepted_value(rsp_msg, pins_impl_.accepted_value());
700         break;
701     }
702
703     assert(false == write);
704     return std::make_tuple(0, write, move(rsp_msg));
705 }
706

```

```

707 void PinsWrapper::markChosen()
708 {
709     pins_impl_.set_chosen(true);
710     if (nullptr != pins_state_) {
711         pins_state_->MarkChosen();
712         assert(pins_state_->IsChosen());
713     }
714 }
715
716 std::tuple<int, bool, std::unique_ptr<Message>>
717 PinsWrapper::stepNotChosen(const Message& msg)
718 {
719     assert(false == IsChosen());
720     if (0 == access(
721         "/home/qspace/data/kvsrv/plog_learner_only", F_OK)) {
722         if (MessageType::CHOSEN != msg.type()) {
723             logerr("plog_learner_only msgtype %d",
724                 static_cast<int>(msg.type()));
725             return std::make_tuple(-50221, false, nullptr);
726         }
727
728         assert(MessageType::CHOSEN == msg.type());
729     }
730
731     bool write = false;
732     MessageType rsp_msg_type = MessageType::NOOP;
733     switch (msg.type()) {
734     // for all
735     case MessageType::NOOP:
736         case MessageType::GET_CHOSEN:
737             // do nothing
738             break;
739
740     case MessageType::CHOSEN:
741         {
742             // FOR NOW
743             assert(true == msg.has_accepted_value());
744             if (pins_impl_.has_accepted_num()
745                 && msg.proposed_num() == pins_impl_.accepted_num()) {
746                 // mark already accepted entry as chosen
747                 assert(pins_impl_.has_accepted_value());
748                 // !! CHECK !!
749                 if ((pins_impl_.accepted_value().reqid() !=
750                     msg.accepted_value().reqid(
751                         pins_impl_.accepted_value().data()
752                         msg.accepted_value().data())) {
753                     logerr("IMPORTANT INCONSISTENT index %" PRI
754                         " from %u to %u",

```



```

755                                     msg.index(), msg.from(), ms
756                                     assert(false);
757                                     }
758                                 }
759                             else {
760                                 // self roll promised, accepted, chosen
761                                 write = true;
762                                 cutils::PropNumGen prop_num_gen(0, 100);
763                                 uint64_t hint_num = std::max(
764                                     msg.proposed_num(), pins_impl_.prom
765                                     hint_num = std::max(hint_num, pins_impl_.proposed_n
766
767                                 logimpt(" index %" PRIu64 " msg: proposed %" PRIu64
768                                     " local: chosen_ %d promised %" PRI
769                                     " hint_num %" PRIu64,
770                                     msg.index(), msg.proposed_num(),
771                                     pins_impl_.chosen(),
772                                     pins_impl_.promised_num(),
773                                     pins_impl_.accepted_num(),
774                                     hint_num);
775
776                                 auto chosen_prop_num = prop_num_gen.Next(hint_num);
777                                 assert(chosen_prop_num > msg.proposed_num());
778                                 assert(chosen_prop_num > pins_impl_.promised_num());
779                                 assert(chosen_prop_num > pins_impl_.proposed_num())
780                                 pins_impl_.set_proposed_num(chosen_prop_num);
781                                 assert(updatePromised(chosen_prop_num, pins_impl_));
782                                 assert(updateAccepted(
783                                     chosen_prop_num,
784                                     msg.accepted_value(), false, pins_impl_));
785                             }
786
787                             markChosen();
788                             // not rsp_msg;
789                         }
790                         break;
791
792                     // acceptor
793                     case MessageType::PROP:
794                         {
795                             if (updatePromised(msg.proposed_num(), pins_impl_)) {
796                                 // promised =>
797                                 write = true;
798                             }
799                             rsp_msg_type = MessageType::PROP_RSP;
800                         }
801                         break;
802

```

```

803     case MessageType::ACCP:
804     case MessageType::FAST_ACCPT:
805         {
806             assert(msg.has_accepted_value());
807             bool fast_accept = MessageType::FAST_ACCPT == msg.type();
808             if (updateAccepted(
809                 msg.proposed_num(),
810                 msg.accepted_value(), fast_accept, pins_impl_) {
811                 // accepted other
812                 write = true;
813             }
814
815             rsp_msg_type = fast_accept
816                 ? MessageType::FAST_ACCPT_RSP : MessageType::ACCP_RSP;
817         }
818         break;
819
820     // proposer
821     case MessageType::PROP_RSP:
822     case MessageType::ACCP_RSP:
823     case MessageType::FAST_ACCPT_RSP:
824         // start a propose
825     case MessageType::BEGIN_PROP:
826     case MessageType::TRY_PROP:
827     case MessageType::BEGIN_FAST_PROP:
828         case MessageType::TRY_REDO_PROP: // new add:
829             if (nullptr == pins_state_) {
830                 logdebug("pins_state nullptr but recv msgtype %d",
831                     static_cast<int>(msg.type()));
832                 break; // just ignore all proposer related msg
833             }
834
835             assert(nullptr != pins_state_);
836             assert(msg.key() == pins_state_->GetKey());
837             assert(msg.index() == pins_state_->GetIndex());
838             std::tie(write, rsp_msg_type) = pins_state_->Step(msg, pins_impl_);
839             break;
840
841     default:
842         assert(false);
843         break;
844     }
845
846     auto rsp_msg = produceRsp(msg, rsp_msg_type);
847     if (MessageType::CHOSEN == rsp_msg_type) {
848         assert(nullptr != rsp_msg);
849         assert(false == IsChosen());
850         markChosen();

```

```

851     }
852     return std::make_tuple(0, write, std::move(rsp_msg));
853 }
854
855 std::unique_ptr<Message>
856 PinsWrapper::produceRsp(const Message& msg, MessageType rsp_msg_type)
857 {
858     if (MessageType::N00P == rsp_msg_type) {
859         return nullptr;
860     }
861
862     std::unique_ptr<Message> rsp_msg = cutils::make_unique<Message>();
863     assert(nullptr != rsp_msg);
864     rsp_msg->set_key(msg.key());
865     rsp_msg->set_index(msg.index());
866     rsp_msg->set_from(msg.to());
867     rsp_msg->set_type(rsp_msg_type);
868     rsp_msg->set_proposed_num(msg.proposed_num());
869     rsp_msg->set_timestamp(time(NULL));
870     switch (rsp_msg_type) {
871
872     // acceptor
873     case MessageType::PROP_RSP:
874         rsp_msg->set_promised_num(pins_impl_.promised_num());
875         assert(rsp_msg->promised_num() >= rsp_msg->proposed_num());
876
877         // TODO: add test
878         // promised or reject => both need send back accepted_num
879         if (pins_impl_.has_accepted_num()) {
880             assert(pins_impl_.has_accepted_value());
881             rsp_msg->set_accepted_num(pins_impl_.accepted_num());
882             set_accepted_value(rsp_msg, pins_impl_.accepted_value());
883         }
884
885         rsp_msg->set_to(msg.from());
886         break;
887
888     case MessageType::ACCP_T_RSP:
889     case MessageType::FAST_ACCPT_RSP:
890         {
891             assert(pins_impl_.has_promised_num());
892             // => reject ?
893             // assert(pins_impl_.has_accepted_num());
894             // assert(pins_impl_.has_accepted_value());
895
896             // TODO: ? send back pins_impl_.promised_num as a hint
897             auto accepted_num =
898                 pins_impl_.has_accepted_num() ? pins_impl_.accepted_num() : 0;

```

```
899     rsp_msg->set_accepted_num(accepted_num);
900     if (accepted_num != msg.proposed_num()) {
901         // reject => return promised_num as a hint
902         if (0 == accepted_num) {
903             rsp_msg->set_promised_num(pins_impl_.promised_num());
904         }
905     }
906     rsp_msg->set_to(msg.from());
907 }
908 break;
909
910 // proposer
911 case MessageType::PROP:
912     assert(nullptr != pins_state_);
913     rsp_msg->set_proposed_num(pins_impl_.proposed_num());
914     // set_to 0 => broad-cast
915     rsp_msg->set_to(0);
916     assert(0 < cutils::get_prop_cnt(rsp_msg->proposed_num()));
917     break;
918
919 case MessageType::ACCP:
920 case MessageType::FAST_ACCP:
921     assert(nullptr != pins_state_);
922     assert(pins_impl_.has_promised_num());
923     assert(pins_impl_.has_accepted_num());
924     assert(pins_impl_.has_accepted_value());
925
926     assert(pins_impl_.proposed_num() == pins_state_->GetProposedNum());
927     assert(pins_state_->HasProposingValue());
928     assert(pins_impl_.accepted_value().data() ==
929            pins_state_->GetProposingValue().data());
930     assert(pins_impl_.accepted_value().reqid() ==
931            pins_state_->GetProposingValue().reqid());
932     rsp_msg->set_proposed_num(pins_state_->GetProposedNum());
933     set_accepted_value(rsp_msg, pins_state_->GetProposingValue());
934     rsp_msg->set_to(0);
935
936     // check
937     if (MessageType::ACCP == rsp_msg_type) {
938         assert(0 < cutils::get_prop_cnt(rsp_msg->proposed_num()));
939     }
940     else {
941         assert(MessageType::FAST_ACCP == rsp_msg_type);
942         assert(0 == cutils::get_prop_cnt(rsp_msg->proposed_num()));
943     }
944     break;
945
946 case MessageType::CHOSEN:
```

```
947     assert(MessageType::CHOSEN != msg.type());
948     assert(pins_impl_.has_promised_num());
949     assert(pins_impl_.has_accepted_num());
950     assert(pins_impl_.has_accepted_value());
951     rsp_msg->set_proposed_num(pins_impl_.proposed_num());
952     set_accepted_value(rsp_msg, pins_impl_.accepted_value());
953     rsp_msg->set_to(0);
954     break;
955 default:
956     assert(false);
957     break;
958 }
959
960 assert(rsp_msg->from() == msg.to());
961 return rsp_msg;
962 }
963
964 } // namespace paxos
965
```